

# EFFECTIVE PLATFORM BUILDING WITH KUBERNETES.



## IS K8S NEW LINUX?

Wojciech Barczynski - [SMACC.io](https://smacc.io) | [Hypatos.ai](https://hypatos.ai)  
1 October 2018

# WOJCIECH BARCZYŃSKI

- Lead Software Engineer  
& System Engineer
- Python i Golang
- Hobby:  
uczenie inżynierii  
oprogramowania



# BACKGROUND

- Machine Learning FinTech
- Before:
  - 1 z 10 Indonesian mobile e-commerce
- Spent 3.5y datacenters z Openstack (0 ->21)
- SAP R&D

# STORY

- Lyke - [12.2016 - 07.2017] - GKE
- SMACC - [10.2017 - present] - OnPrem + AKS

# NIE LUBIE INFRA :D


- Fabric, Ansible
- Puppet... Chef
- ...
- CloudFormations...

Terraform <3

# **NIE LUBIE INFRA :D**

- Administracja jest trudna i kosztowna
- Virtualne Maszyny, ansible, salt, etc.
- Za dużo ruchomych części
- Nie kończąca się standaryzacja

# MIKROSERWISY AAA!



I...

- Cloud is not so cheap - \$\$\$
- Cloud IaaS lock-in is real



# A GDYBYŚ

- nie musiał myśleć o IaaS
- nigdy więcej logowania się na VM
- mniej dyskusji o CI / CD
- DC jako czarna skrzynka


# KUBERNETES

- Container management
- Service and application mindset
- Simple Semantic\*
- Independent from IaaS provider

# KUBERNETES


- Batteries for your 12factory apps
- Service discovery, meta-data support
- Utilize resources to nearly 100%

# KUBERNETES




`make docker_push; kubectl create -f app-srv-dpl.yaml`

# SCALE UP! SCALE DOWN!




```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

# SCALE UP! SCALE DOWN!



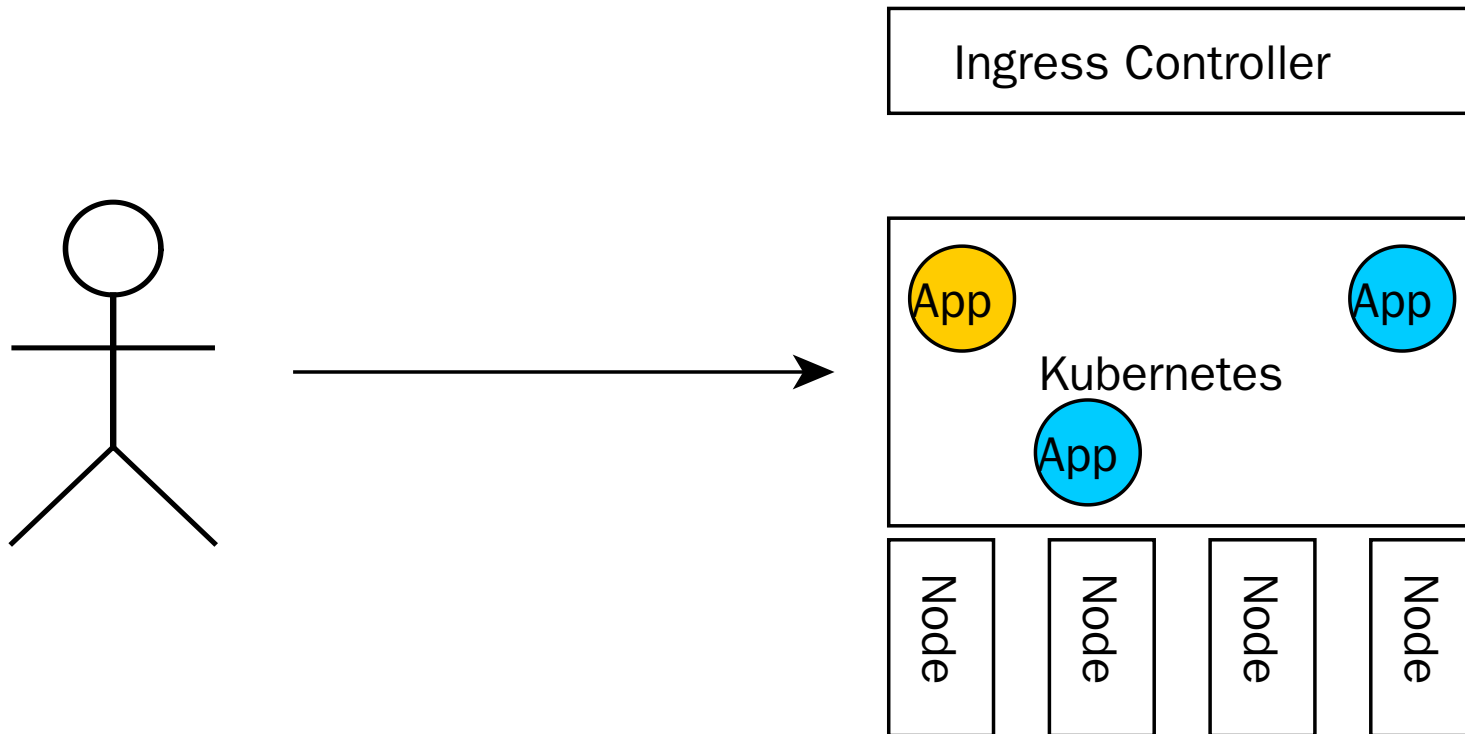
```
kubectl --replicas=1 -f app-srv-dpl.yaml
```

# ROLLING UPDATES!



```
kubectl set image deployment/app app=app:v2.0.0
```


# ROLLING UPDATES!



```
kubectl set image deployment/app app=app:v2.0.0
```




# ROLLING UPDATES!




```
kubectl set image deployment/app app=app:v2.0.0
```


# RESISTANCE!



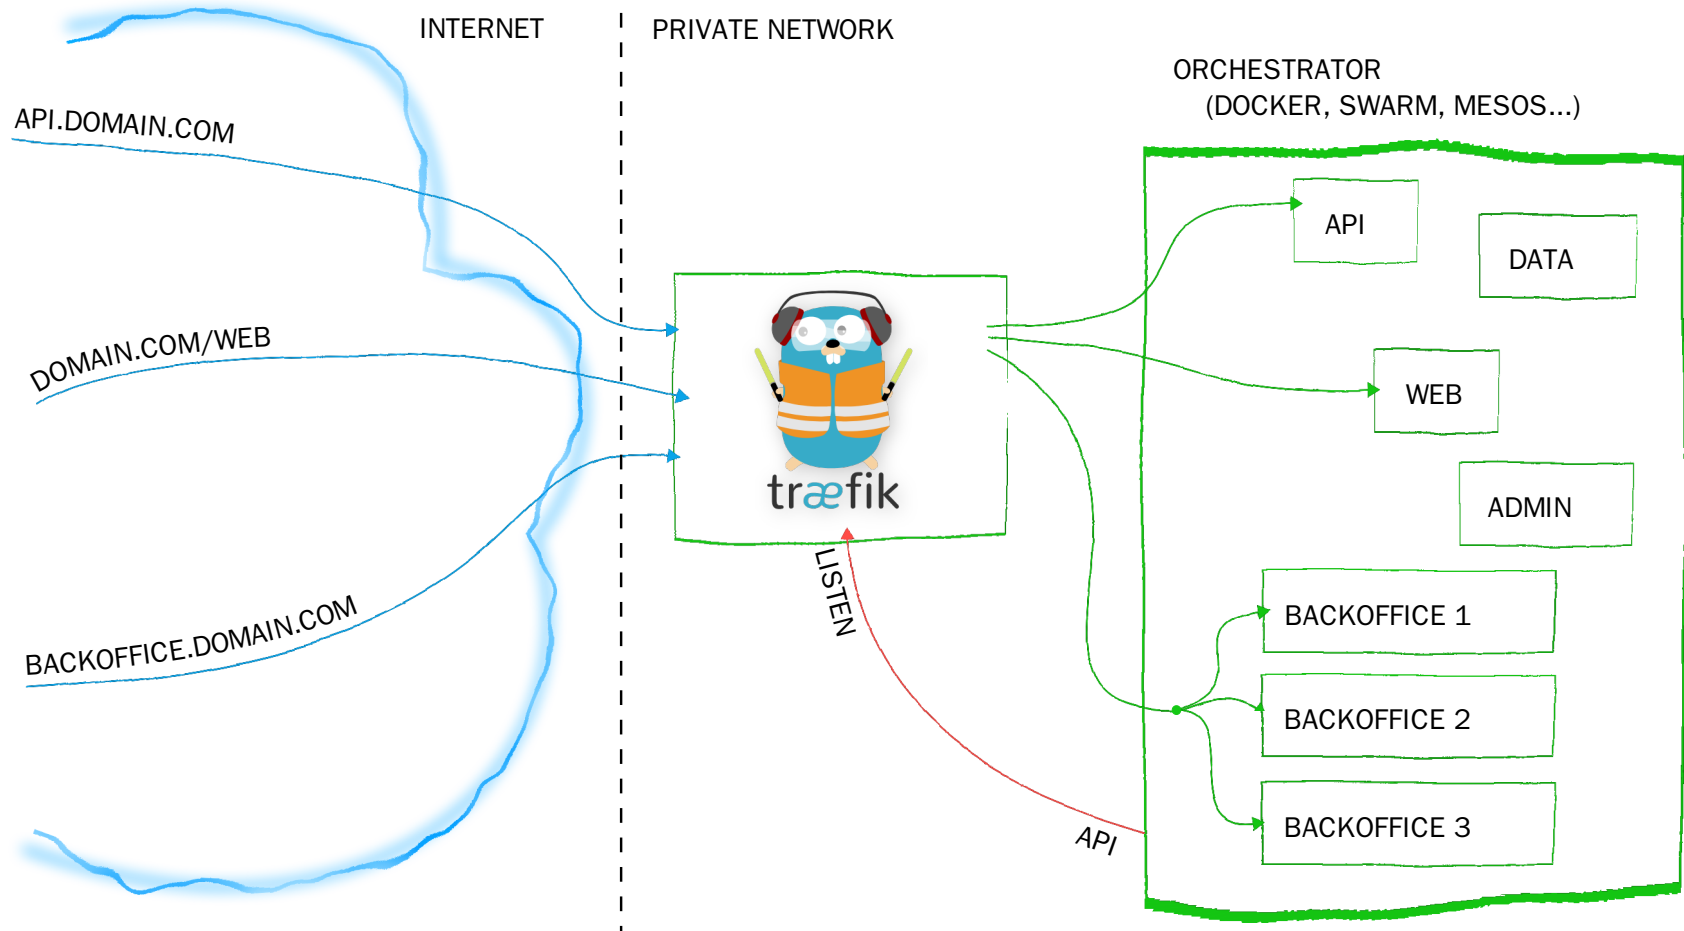
# RESISTANCE!



# RESISTANCE!



# HOW GET USER REQUESTS?




Ingress Controller

# INGRESS

Pattern	Target App Service
api.smacc.io/v1/users	users-v1
api.smacc.io/v2/users	users-v2
smacc.io	web

# LOAD BALANCING



# SERVICE DISCOVERY

- names in DNS:

```
curl http://users/list
```

- labels:

```
name=value
```

- annotations:

```
prometheus.io/scrape: "true"
```



# AUTO-WIRING

```
---
apiVersion: v1
kind: Service
metadata:
  name: memcached-0
  labels:
    app: memcached
    role: shard-0
    tier: backend
  annotations:
    prometheus.io/scrape: "true"
    prometheus.io/scheme: "http"
    prometheus.io/path: "metrics"
    prometheus.io/port: "9150"
```

<https://github.com/skarab7/kubernetes-memcached>

# AUTO-WIRING

```
groups:
- name: apps
  rules:
- alert: ProductionInstanceDown
  expr: up{env = 'production'} == 0
  for: 2m
  labels:
    severity: opsgenie
  annotations:
    summary: "Instance {{ $labels.instance }} of {{ $labels."
    description: "{{ $labels.instance }} of job {{ $labels.j
```

# DROP-IN

- traefik / Ingress / Envoy
- prometheus
- audit checks
- ...

# MORE

- readiness probe
- liveness probe
- resource quotas

# ALL IN GIT

- all in yaml
- integration with monitoring, alarming
- integration with ingress-controller
- ...
- devs can forget about infrastructure... almost

DevOps Culture Dream!

I..


- łatwo zaimplementować Continuous Deployment
- i ukryć detale



LYKE

# LYKE

- Mobile E-commerce
- 50k+ użytkowników
- 2M downloads
- Top 10 Fashion Apps



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Teraz: JollyChic Indonesia



# GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,  
UI Designer,  
Mobile Dev,  
and tester - one body



## CHALLENGES

- 50+ VMs in Amazon
- 1 VM - 1 App, idle machine
- Puppet, hilarious (manual) deployment process
- Fear
- Forgotten components
- Performance issues

# APPROACH

1. Simplify infrastructure
2. Change the Development practices
3. Change the work organization

see: Conway's law

## **SIMPLIFY**

1. Kubernetes with Google Kubernetes Engine
2. Terraform for all new

## SIMPLIFY

1. Monitoring: Prometheus + Grafana
2. Logging: Elasticsearch-Fluentd-Kibana
3. Google Identity-Aware-Proxy to protect tools
4. 3rd party SaaS: statuscake and opsgenie

# CONTINUOUS DEPLOYMENT

- branch-based:
  - master
  - staging
  - production
- repo independent

# TRAVIS CI

1. Tests
2. Build docker
3. Deploy to Google Container Registry
4. Deploy to k8s only new docker
5. no config applied

# GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \ - kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```



# Makefile

```
SERVICE_NAME=v-connector
GCP_DOCKER_REGISTRY=eu.gcr.io
test: test_short test_integration

run_local:

docker_build: docker_push

kube_create_config:

kube_apply:

kube_deploy:
```

Copy&Paste from the project to project

# 1. CLEAN UP

- Single script for repo - Makefile [1]
- Resurrect the README

[1] With zsh or bash auto-completion plug-in in your terminal.

## 2. GET BACK ALL THE KNOWLEDGE

- Puppet ... ➡ Dockerfile
- Check the instances ➡ Dockerfile, README.rst
- Nagios, ... ➡ README.rst, [checks/](#)

### 3. INTRODUCE RUN\_LOCAL

- `make run_local`
- A nice section on how to run in README.rst
- Use: `docker-compose`

The most crucial point.

## 4. GET TO KUBERNETES

- `make kube_create_config`
- `make kube_apply`
- Generate the yaml files if your envs differ

## 5. CONTINUOUS DEPLOYMENT

Travis:

- use the same Makefile as a developer


## 6. KEEP IT RUNNING

Bridge the new with old:

- Use external services in Kubernetes
- Optional: Expose k8s in the Legacy [1]

[1] feeding K8S events to HashiCorp consul


# Bridge the new with old



Monitor legacy with new stack



# Architecture During Migration



## 7. INTRODUCE SMOKE-TEST

```
TARGET_URL=127.0.0 make smoke_test
```

```
TARGET_URL=api.example.com/users make smoke_test
```

## 8. MOVE TO MICRO-SERVICES

To offload the biggest components:

- Keep the lights on
- New functionality delegated to micro-services

## 9. SERVICE SELF-CONSCIOUSNESS

Add to old services:

1. *metrics/*
2. *health/*
3. *info/*

## 10. GET PERFORMANCE TESTING

- introduce *wrk* for evaluating performance
- load test the real system

# WHAT WORKED

1. Copy&Paste Makefile and k8s
2. Separate deployments a good transition strategy

# WHAT DID NOT WORK

1. Too many PoC ➡ 2 weeks max
2. Do it with smaller chunks
3. Alert rules too hard to write
4. Push-back to k8s yaml [\*]

[\*] With coaching, I thought, it is OK

# DO DIFFERENT

1. Move dev and staging data immediately
2. Let devs know it is a transition stage
3. Teach earlier about quotas
4. EFK could wait
5. Big migration - a paid-XXX% weekend

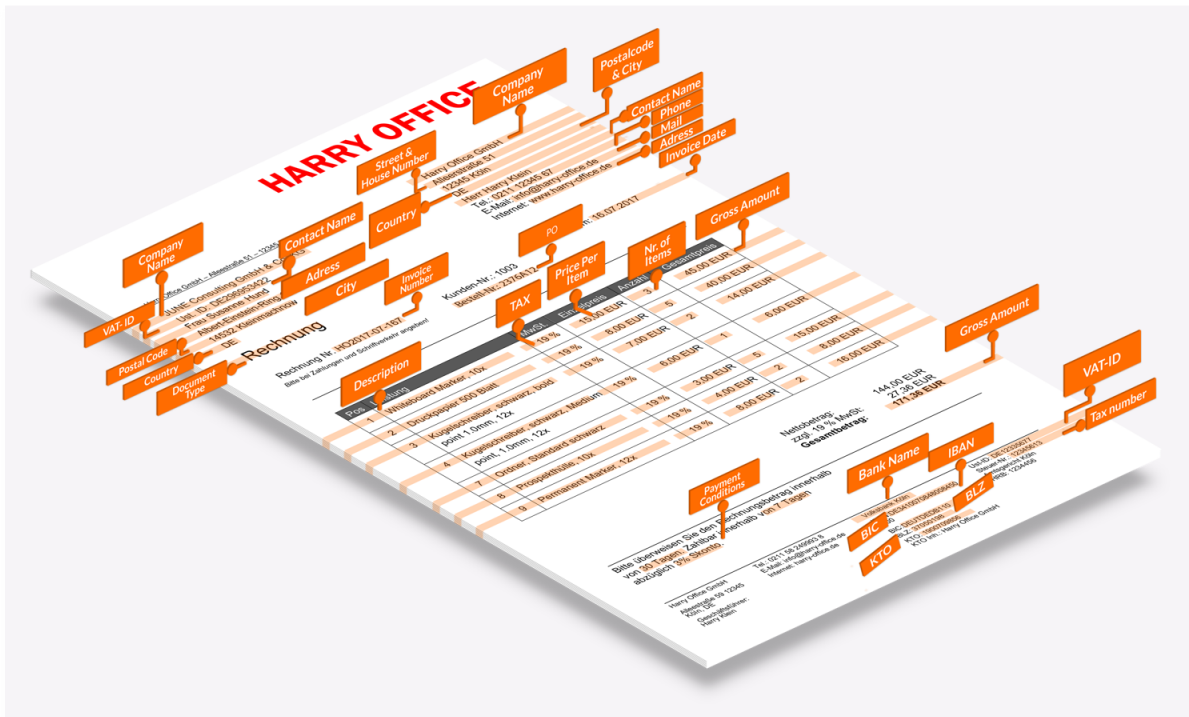


SMACC



Hypatos

# SMACC.io



SMACC

# Klienci

Deutsche Bank



BMB  
巴三巴  
GRUPPE



Deloitte.



>  
**accenture**

McKinsey&Company

SMACC

# STORY

- Legacy on AWS
- Experiments with AWS ECS :/
- Self-hosted K8S on ProfitBricks
- Get to Microsoft ScaleUp, welcome Azure!

# LUCKILY - AKS



Azure



Nie było, aż tak różowo ☠

# AZURE KUBERNETES SERVICE

- Niezależne od IaaS
- Our OnPrem = Our OnCloud
- Konsolidacja naszych mikroservisów
- Plug & play, e.g., monitoring

# PROSTOTA

- az aks CLI do ustawienia k8s - README.rst
- Terraform dla wszystkiego wokół
- 1Password i gopass.pw

TF dla AWS

# RÓŻNICA

- Dwa zespoły: Berlin i Warszawa
- Ja w Warszawie



# NOWE DOŚWIADCZENIE

- devs really do not like TravisCI ... k8s yamls
- ciężkie przejście z ProfitBricks do AKS

# ROZWIĄZANIE

- prościej :)
- copy&paste dla wszystkiego
- ukryć deployment na k8s, usunąć magic
- deployment on *tag*

Similar to the [Kelsey Hightower approach](#)

# Repo .travis.yml

```
language: go
go:
- '1.10'
services:
- docker
install:
- curl -sL https://${GITHUB_TOKEN}@raw.githubusercontent.com
- if [ -f "tools/travis/install.sh" ]; then bash tools/travi
script:
- dep ensure
- make lint
- make test
- if [ -z "${TRAVIS_TAG}" ]; then make snapshot; fi;
deploy:
- provider: script
```

# Makefile

```
| - tools  
|   | - Makefile  
|   | - kube-service.yaml  
|   \- kube-deployment.yaml  
|  
| - Dockerfile  
\- Makefile
```

# CONTINUOUS DEPLOYMENT

- Github
- TravisCI
- [hub.docker.com](https://hub.docker.com)
- AKS

# PROCESS

1. `git tag` and push
2. Generate deploy, ingress, and svc kubernetes files
3. Merge PR -> production

# KUBERNETES

- Pure, generated, kubernetes config
- 2x kubernetes operators

# NEXT

- Acceptance tests
- Scale our Machine Learnings
- Deployment tool based on [missy](#)
- Keeping an eye on Istio



# Kubernetes - Linux

- Not a silver bullet, but damn close
- Common runtime dla onPrem i onCloud
- Z kubevirt - zastąpić może Openstack

# Kubernetes - Linux

- The biggest asset - the API
- With service discovery - an integration platform

DZIĘKUJĘ. PYTANIA?

ps. We are hiring.

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE  
BE WITH YOU.




# BACKUP SLIDES


```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

MAY THE SOURCE  
BE WITH YOU.




# KUBERNETES CONCEPTS







# PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports





# SIDE-CARS



## BASIC CONCEPTS

**Name**

**Purpose**

Service

Interface

Entry point  
(Service Name)

---

Deployment

Factory

How many pods,  
which pods

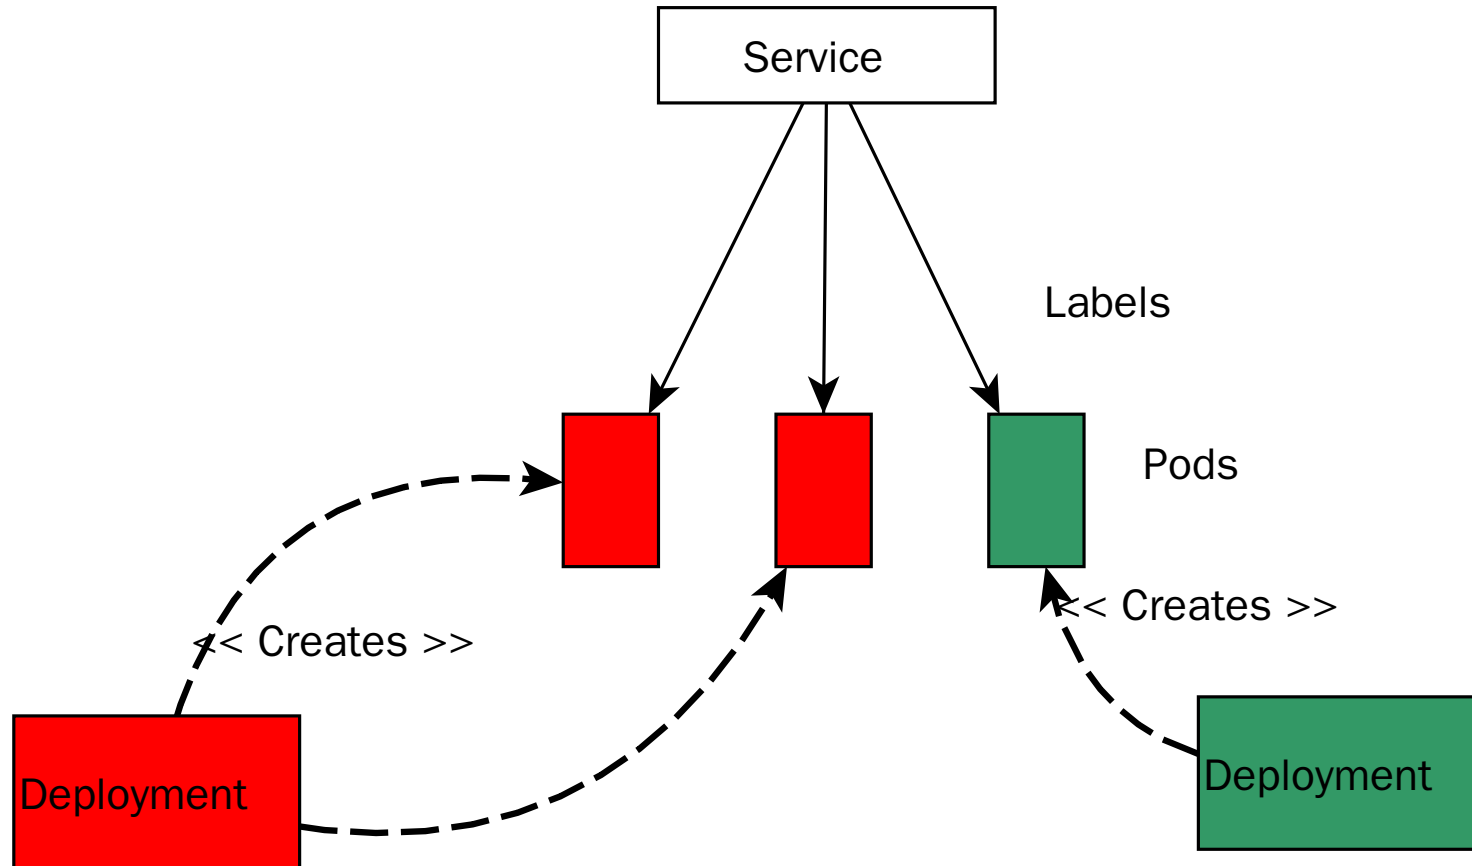
---

Pod

Implementation

1+ docker running

# ROLLING RELEASE WITH DEPLOYMENTS



Also possible

# KUBERNETES CONCEPTS

# MACHINE LEARNING WORKLOADS

- Rediness probe
- Liveness probe

# K8S CLIENTS

- golang
- python